

1 Computational Mechanics Meets Artificial Intelligence

Leon Herrmann

Stefan Kollmannsberger

Chair of Data Engineering in Construction

Bauhaus-Universität Weimar

Paris, February 2025

*Deep Learning in Computational Mechanics – an introductory course,
Herrmann et al. 2025*



website



book



Contents

- 1 Computational Mechanics Meets Artificial Intelligence (& Introduction to PyTorch):
 - What is Artificial Intelligence?
 - History of Artificial Intelligence
 - Recent Achievements of Artificial Intelligence
 - Artificial Intelligence in Science
 - Challenges
 - Computational Mechanics Meets Artificial Intelligence
- 2 Fundamental Concepts of Machine Learning
- 3 Neural Networks
- 4 Introduction to Physics-Informed Neural Networks
- 5 Advanced Physics-Informed Neural Networks
- 6 Machine Learning in Computational Mechanics
- 7 Material Modeling with Neural Networks
- 8 Generative Artificial Intelligence
- 9 Inverse Problems & Deep Learning
- 10 Methodological Overview of Deep Learning in Computational Mechanics
- 11 The Future of Deep Learning in Computational Mechanics

What is Artificial Intelligence?

Artificial Intelligence: A Modern Approach, Norvig et al. 2020

Artificial Intelligence

- “Intelligence exhibited by machines/computers”
- (Total) **Turing test** requires: natural language processing, knowledge representation, automated reasoning, machine learning, (computer vision, robotics)

Intelligence


- **Human** or **rational**?
- **Intelligent thoughts** or **intelligent behavior**?

Machine Learning

- “Learn from data & generalize to unseen data (without explicit instructions)”

Deep Learning

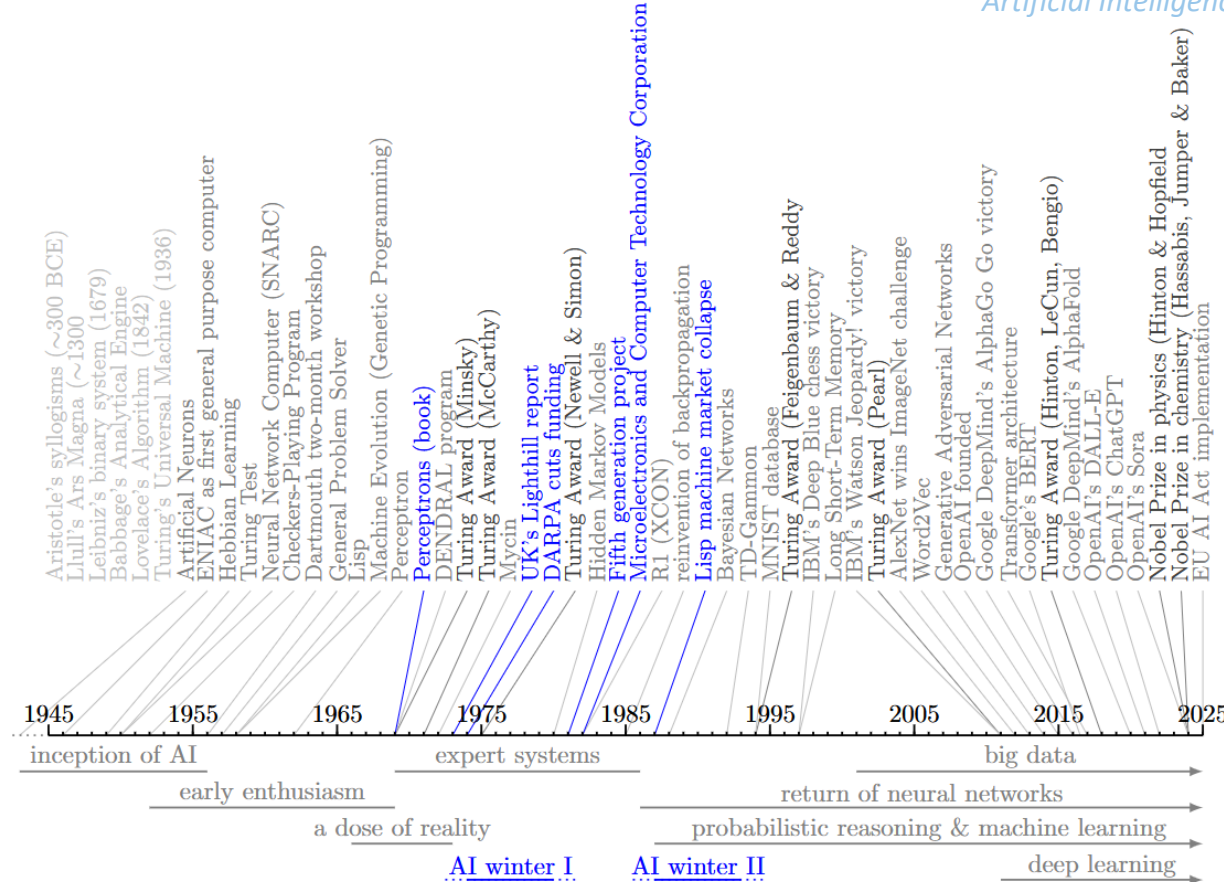
- “Training (deep) **neural networks**”



Inspired by Rebekka Woldseth, author of "On the use of artificial neural networks in topology optimisation"

History of Artificial Intelligence

Artificial Intelligence: A Modern Approach, Norvig et al. 2020



Superintelligence: Paths, Dangers, Strategies, Bostrom 2014

Artificial General
Intelligence?
⋮
AI winter III?

History of Artificial Intelligence

Artificial Intelligence: A Modern Approach, Norvig et al. 2020

- **The inception of artificial intelligence (1943-1956)**
 - Basic physiology of the brain → **artificial neurons** (on/off); updating rule as Hebbian Learning; SNARC
- **Early enthusiasm, great expectations (1952-1969)**
 - Turing “a machine can never do X”; models were based on logic and symbolic reasoning; (GPS, Lisp, perceptron)
- **A dose of reality (1966-1973)**
 - Overconfidence: models based on “informed introspection” & “intractability of attempted problems”; Lighthill
- **Expert systems (1969-1986)**
 - Instead of general-purpose tools; **domain-specific knowledge**; (DENDRAL, Mycin, R1); Fifth Generation Project
- **The return of neural networks (1986-)**
 - Reinvention of **backpropagation**
- **Probabilistic reasoning and machine learning (1987-)**
 - Reaction to failure of expert systems; **learn from experience** → adaptable & incorporation of uncertainty
 - Hidden Markov Models (Reinforcement Learning); Bayesian Networks; TD-Gammon
- **Big data (2001-)**
 - World Wide Web: **large datasets** (billions-trillions of samples); **ImageNet** (challenge), IBM’s Watson
- **Deep learning (2011-)**
 - Hardware improvements (**GPU**: $10^{14} - 10^{17}$ vs CPU: $10^9 - 10^{10}$ Flops); (**Deep CNNs** in AlexNet); AlphaGo

Recent Achievements in Artificial Intelligence




<https://media.freemalaysiatoday.com/wp-content/uploads/2022/05/lifestyle-garry-emel-pic-110522.jpg>




Cats versus dogs



<https://media.freemalaysiatoday.com/wp-content/uploads/2016/03/AlphaGo.jpg>



https://commons.wikimedia.org/wiki/File:C12orf29_AlphaFold.png



What can I help with?

Message ChatGPT



Create image

Analyze data


Summarize text

Get advice

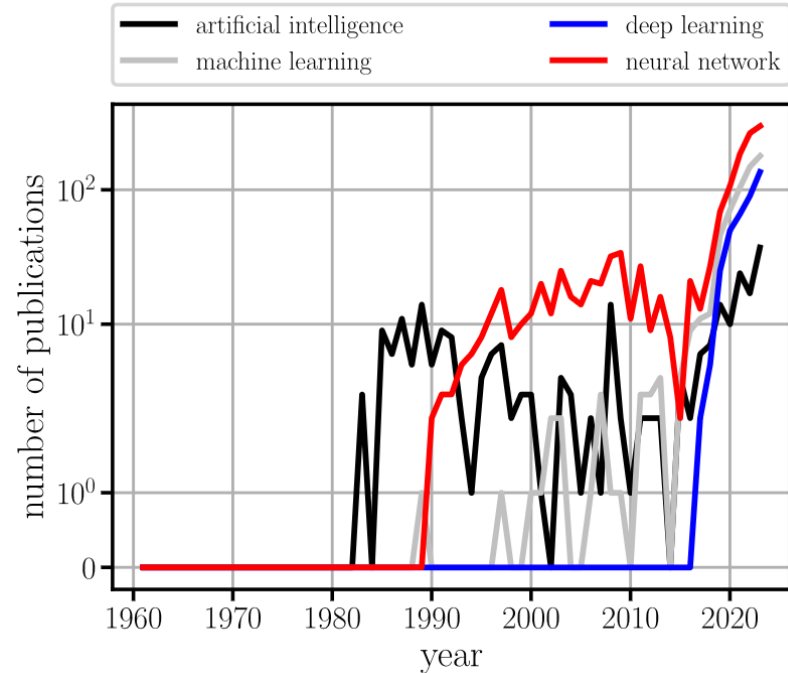
More

Artificial Intelligence in Science

Check www.aitracker.org for other trends




Publications in all fields



Publications in computational mechanics


Challenges

- “Generate an image of Isaac Newton in front of a blackboard on which his three laws are written in mathematical notation and chalk.”
- Follow-up: “The laws on the blackboard are incorrect. Please add the correct formulations. If you are unable to do so, simply focus on the second law, which is $F=m \cdot a$.”



Generated with
DALL-E-3

Challenges



<https://openai.com/index/attacking-machine-learning-with-adversarial-examples/>

Limitations in deep learning in general

- Neural networks **break** in unpredictable ways → can be consistently fooled
- Deep learning is **not robust** due to sensitivity to **hyperparameters** → requires extensive tuning
- Neural networks are **uninterpretable**, i.e., limited explainability → limits reliability

See chapter 11 for details

Problems in deep learning in computational mechanics

- **Reproducibility crisis** (bias towards positive results, sensitivity, transparency)
- **Fair evaluation metrics** are disregarded (breakeven threshold, meaningful metrics, statistical assessments)
- **State-of-the-art** is not considered

$$\tau = \frac{T_{\text{data}} + T_{\text{train}}}{T_{\text{simulation}} - T_{\text{surrogate}}}$$

Good scientific practice for deep learning in computational mechanics

- Honest assessments & explanations (consider the **state-of-the-art & proper metrics**)
- Proposed methods should be **robust** towards **hyperparameters** (no extensive tuning for a novel problem)
- Careful & **narrower selection** of problem types (**not general-purpose** solution)
 - domain-specific improvements

Towards a meaningful integration of neural networks in computational solid mechanics, Herrmann 2025

Example from topology optimization

The mean squared error

$$MSE = \frac{1}{m} \sum_{i=1}^m (x_{\text{left}_i} - x_{\text{right}_i})^2$$


between the two structures is very small ($2.5 \cdot 10^{-3}$), due to one pixel difference.

Structural compliance (inverse of stiffness)


$$c = \mathbf{F}^T \mathbf{u}$$

Is different by one order of magnitude

For more details,
see Chapter 9



compliance= 351



compliance= 4729

Computational Mechanics Meets Artificial Intelligence

Computational Mechanics

Abstraction of physical systems (reality) through simplified mathematical models (often differential equations), which are discretized and solved numerically for insight into real-world behavior

Exemplary tasks

- Efficient solutions techniques for **forward problems**, e.g., finite element, difference, and volume
- Identification tasks (**inverse problems**), e.g., inferring material distribution/properties from measurements
- **Optimization**, e.g., finding the optimal material distribution that maximizes stiffness

Machine Learning

Machine Learning, Mitchell 1997

"a computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks T , as measured by P , improves with experience E "

Where can machine learning be applied in computational mechanics?

- **Identification** of mathematical models from data (instead of relying on hand-crafted models)
- **Acceleration** of forward solvers and optimizers
- **Streamlining** of pipelines to avoid human experts within the processes

Computational Mechanics Meets Artificial Intelligence

Deep learning in computational mechanics: a review, Herrmann et al. 2024

- **Simulation substitution**

- Data-driven modelling
- Physics-informed learning

Simulation with graph neural networks; DMD; Transfer learning

Hamiltonian/Lagrangian neural networks; SINDy; (PINNs)

- **Simulation enhancement**

Input-convex neural networks for material modeling; EUCLID; Neural networks as ansatz function of inverse quantities; Super resolution; Differentiable physics

- **Discretizations as neural networks**

Hardware acceleration with GPUs; (HiDeNN)

- **Generative approaches**

Generative design; Realistic data generation; Anomaly detection; Transformers for natural language processing

- **Deep reinforcement learning**

Control engineering tasks: autonomous flight; robots; Alternative gradient-free optimizer

2 Fundamental Concepts of Machine Learning

Leon Herrmann

Stefan Kollmannsberger

Chair of Data Engineering in Construction

Bauhaus-Universität Weimar

Paris, February 2025

*Deep Learning in Computational Mechanics – an introductory course,
Herrmann et al. 2025*



website



book



Contents

- 1 Computational Mechanics Meets Artificial Intelligence (& Introduction to PyTorch)
- 2.1 Definition
- 2.2 Data Structure
- 2.3 Types of Learning
- 2.4 Machine Learning Tasks
- 2.5 Linear Regression
- 2.8.1 Gradient Descent
- 3 Neural Networks


2.1 Definition

“a computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks T , as measured by P , improves with experience E ”

Machine Learning, Mitchell 1997

Most machine learning algorithms are composed of

- Dataset
- Parametrized model
- Cost function
- Optimization procedure



2.2 Data Structure

Specific dataset (sometimes called design matrix)

$$\mathbf{X} = \begin{matrix} & \text{feature 1} & \text{feature 2} & \cdots & \text{feature } n \\ \text{example 1} & x_{11} & x_{12} & \cdots & x_{1n} \\ \text{example 2} & x_{21} & x_{22} & \cdots & x_{2n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \text{example } m & x_{m1} & x_{m2} & \cdots & x_{mn} \end{matrix}$$

Examples

- Can be different images, where its features are its pixel values, $n = \text{channels} \times \text{pixels}$
- Can be different houses, where its features are its properties such as area, number of rooms, age

Notation

- Design matrix \mathbf{X}
- Design vector of a single example i (1 sample/example) $\mathbf{x}_i = [x_{i1}, x_{i2}, \dots, x_{in}]^T$

2.3 Types of Learning

$$\mathbf{X} = \begin{matrix} & \text{feature 1} & \text{feature 2} & \cdots & \text{feature } n \\ \text{example 1} & x_{11} & x_{12} & \cdots & x_{1n} \\ \text{example 2} & x_{21} & x_{22} & \cdots & x_{2n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \text{example } m & x_{m1} & x_{m2} & \cdots & x_{mn} \end{matrix}$$

Supervised learning

- Algorithm learns from a **labeled dataset**. Each **sample** \mathbf{X}_i has an accompanying **target** y_i
- Example: A model learns to distinguish between dogs and cats via annotated images

Unsupervised learning

- Algorithm finds a structure or **pattern** in the data. This is typically in the form of a probability distribution
- Example: Anomaly detection, i.e., the identification of irregularities in otherwise regular patterns. For example in the detection of tumors in medical imaging

Semi-supervised learning

- Combination of supervised and unsupervised learning, i.e., the data is partly labeled to improve the unsupervised learning.
- Example: The learning of the tumor identification is improved by using some labeled data.

Reinforcement learning

- Interaction between an algorithm and an **environment**, improving the algorithm to **maximize** an expected average **reward**. Common in game-like environments
- Example: The stock market, where more actions with higher rewards are learned

2.4 Machine Learning Tasks


Regression

- Prediction of a numerical value via a ([real-valued](#)) mapping between input and output
- Example: Prediction of house prices from criteria like area, number of rooms, age

Classification

- Prediction of a discrete category via a mapping between input and a ([discrete](#)) category
- Example: Classification of images in cats and dogs

Classification can be regarded as discrete regression.



Regression


2.4 Machine Learning Tasks

Clustering

- Discovers similarities between data and creates discrete clusters (unsupervised)
- Example: Identification of similar customer groups

Generative modeling

- Generate new data points that resemble a given dataset (without simply reproducing given data points)
- Example: Generate new rim designs given a set of rims



Clustering

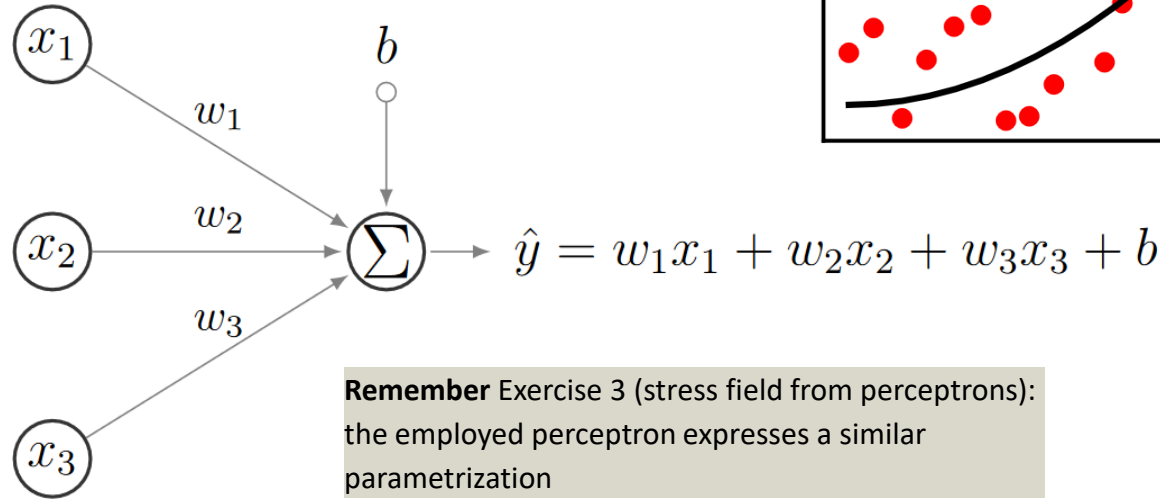
Machine Learning Algorithms

Machine Learning in Additive Manufacturing: State-of-the-Art and Perspectives, Wang et al. 2020



	Classifications	Algorithms	Tasks
	Supervised	Decision trees	Classification
		Random forest	Classification, regression
covered in Chapter 7		Support vector machines	Classification, regression
		K-nearest neighbours	Classification
		Bayesian network	Classification
		Gaussian process	Regression
		Multi-gene genetic programming	Regression
		Hidden semi-Markov model	Classification
covered in Chapter 3		Multi-layer perceptron	Classification, regression
		Convolutional neural network	Classification
		Recurrent neural network	Time series prediction (regression)
		Adaptive network-based fuzzy inference system	Regression
covered in Chapter 8		Transformers	Regression, classification, generative modeling
	Unsupervised	Self-organizing map	Clustering
		Deep belief network	Classification
covered in Chapter 7		K-means clustering	Clustering
		Reduced order modeling (POD)	Dimensionality reduction
covered in Chapter 8		Autoencoder	Generative modeling, dimensionality reduction
		Generative adversarial networks	Generative modeling, (classification)
		Diffusion model	Generative modeling
	Semi-supervised	Gaussian mixture model	Clustering

2.5 Linear Regression – Prediction

- Target: \hat{y}
- Ground truth: \tilde{y} (associated with $\tilde{\mathbf{x}}$)
- Example vector: \mathbf{x}
- Weight vector: \mathbf{w}
- Bias: b



Remember Exercise 3 (stress field from perceptrons):
the employed perceptron expresses a similar
parametrization



2.5 Linear Regression – Performance Measurement

Prediction in nd for sample i :

$$\hat{y}_i = \mathbf{w} \cdot \mathbf{x}_i + b = \sum_{j=1}^n w_j x_{ij} + b$$

i is an example/sample and j is a feature

- each feature has an associated weight, which is shared across all samples

Remember the design matrix

$$\mathbf{X} = \begin{matrix} & \text{feature 1} & \text{feature 2} & \cdots & \text{feature } n \\ \text{example 1} & x_{11} & x_{12} & \cdots & x_{1n} \\ \text{example 2} & x_{21} & x_{22} & \cdots & x_{2n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \text{example } m & x_{m1} & x_{m2} & \cdots & x_{mn} \end{matrix}$$

Where each example vector is defined as $\mathbf{x}_i = (x_{i1}, x_{i2}, \dots, x_{in})^T$

2.5 Linear Regression – Performance Measurement

Prediction in semi-vector notation for sample i

$$\hat{y}_i = \mathbf{w} \cdot \tilde{\mathbf{x}}_i + b$$

Squared error for a single sample i (with **ground truth** \tilde{y}_i)

$$(\tilde{y}_i - \hat{y}_i)^2$$

Mean squared error (MSE) for a dataset X with m samples (including the design vectors \mathbf{x}_i of each sample i)

$$\mathcal{L} = \frac{1}{m} \sum_{i=1}^m (\tilde{y}_i - \hat{y}_i)^2 = \frac{1}{m} \sum_{i=1}^m (\tilde{y}_i - \mathbf{w} \cdot \tilde{\mathbf{x}}_i + b)^2$$

Cost function

$$C(\mathbf{w}, b) = \mathcal{L} + \dots$$

Optimization problem

$$\min_{\mathbf{w}, b} C(\mathbf{w}, b) = \min_{\mathbf{w}, b} \frac{1}{m} \sum_{i=1}^m (\tilde{y}_i - (\mathbf{w} \cdot \tilde{\mathbf{x}}_i + b))^2$$

In machine learning optimization is referred to as **learning** when the model is applied to previously unseen problems (i.e., datapoints). This stands in contrast to structural optimization in which one specific design is obtained through optimization.

2.5 Linear Regression – Data Split



Data is split into

- **Training set** (~80%)
 - to train the model (i.e., find the correct weights and bias)
- **Validation set** (~10%)
 - to validate the training (i.e., evaluate if training is successful, e.g., to detect/avoid **overfitting**)
 - to find the correct (machine learning algorithm) **hyperparameters**
- **Testing set** (~10%)
 - to test/assess the validity of the final model (i.e., weights, bias, and hyperparameters)
 - At this point nothing is allowed to be changed (otherwise testing set becomes validation set)
 - In AI double-blinded challenges are common (test set is released only after handing model to jury)


2.5 Linear Regression – Data Split

Example

- Fitting of datapoints with a polynomial where the hyperparameter is the polynomial degree p



$p = 3$



$p = 8$

- The validation set shows that $p = 8$ leads to the lowest validation error for the trained model
- The test set shows that some overfitting happened during hyperparameter tuning (this information is not available during/for model development)
- The best model would rely on $p = 3$

2.5 Linear Regression – Optimization

$$\min_{\mathbf{w}, b} C(\mathbf{w}, b) = \min_{\mathbf{w}, b} \frac{1}{m} \sum_{i=1}^m (\tilde{y}_i - (\mathbf{w} \cdot \tilde{\mathbf{x}}_i + b))^2$$

Note that \mathbf{X} requires a column of ones for the bias b

For a more concise notation let us denote all learnable parameters in a vector $\Theta = (\mathbf{w}, b)^T$. This allows to write the model function $\hat{y}_i = \mathbf{w}^T \tilde{\mathbf{x}}_i + b$ as $\hat{\mathbf{y}} = \mathbf{X}\Theta$ yielding the minimization

All predictions \hat{y}_i are collected in the vector $\hat{\mathbf{y}}$.

$$\min_{\Theta} C(\Theta) = \min_{\Theta} (\tilde{\mathbf{y}} - \mathbf{X}\Theta)^T (\tilde{\mathbf{y}} - \mathbf{X}\Theta) = \min_{\Theta} (\tilde{\mathbf{y}}^T \tilde{\mathbf{y}} - 2\tilde{\mathbf{y}}^T \mathbf{X}\Theta + (\mathbf{X}\Theta)^T \mathbf{X}\Theta)$$


The minimization is solved by setting the first derivative of C with respect to Θ to zero (using $\mathbf{r} = \tilde{\mathbf{y}} - \mathbf{X}\Theta$)

$$\frac{1}{2} \frac{\partial r(\Theta)^2}{\partial \Theta} = \frac{1}{2} (-2\mathbf{X}^T \tilde{\mathbf{y}} + 2\mathbf{X}^T \mathbf{X}\Theta) = -\mathbf{X}^T \tilde{\mathbf{y}} + \mathbf{X}^T \mathbf{X}\Theta = 0$$

$$\mathbf{X}^T \mathbf{X}\Theta = \mathbf{X}^T \tilde{\mathbf{y}}$$

$$\Theta = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \tilde{\mathbf{y}}$$

Such a closed form solution is only possible if $\hat{\mathbf{y}}$ (or rather $\partial C / \partial \Theta$) is linear with respect to Θ



2.8.1 Gradient Descent

Improve prediction $\hat{y}_i = \mathbf{w} \cdot \mathbf{x}_i + b$ via (iterative) **cost function minimization**

$$\min_{\mathbf{w}, b} C(\mathbf{w}, b) = \min_{\mathbf{w}, b} \frac{1}{m} \sum_{i=1}^m (\tilde{\mathbf{y}}_i - (\mathbf{w} \cdot \mathbf{x}_i + b))^2$$

Partial derivatives of cost function with respect to each parameter

$$\frac{\partial C}{\partial \mathbf{w}} = \frac{1}{m} \sum_{i=1}^m -2x_i(\tilde{\mathbf{y}}_i - (\mathbf{w} \cdot \mathbf{x}_i + b))$$


$$\frac{\partial C}{\partial b} = \frac{1}{m} \sum_{i=1}^m -2(\tilde{\mathbf{y}}_i - (\mathbf{w} \cdot \mathbf{x}_i + b))$$

Each **gradient descent iteration** updates the parameters, such that the cost function decreases

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \frac{\partial C}{\partial \mathbf{w}}$$

$$b \leftarrow b - \alpha \frac{\partial C}{\partial b}$$

α (the **learning rate**) controls the **step size**



2.8.1 Gradient Descent

- Generalized gradient descent algorithm
- In machine learning:
 - Number of iterations is called **number of epochs**
 - Step size is called **learning rate**

Algorithm 1 Gradient descent

Require: dataset \tilde{x}, \tilde{y} , number of epochs n , step size α , model f
initialize the model $f(x; \Theta)$


for all n **do**

 Compute the cost function $C(f(\tilde{x}; \Theta), \tilde{y})$

 Compute the gradient $\nabla_{\Theta} C$

 Update the model parameters $\Theta \leftarrow \Theta - \alpha \nabla_{\Theta} C$

end for



Exercises

- E.4 Linear Regression (P & C)
 - Perform a linear regression once by computing the weights directly and once using gradient descent. Do this by hand calculation and with a Python implementation.

Contents

- 1 Computational Mechanics Meets Artificial Intelligence (& Introduction to PyTorch):
 - What is Artificial Intelligence?
 - History of Artificial Intelligence
 - Recent Achievements of Artificial Intelligence
 - Artificial Intelligence in Science
 - Challenges
 - Computational Mechanics Meets Artificial Intelligence
- 2 Fundamental Concepts of Machine Learning
 - 2.1 Definition
 - 2.2 Data Structure
 - 2.3 Types of Learning
 - 2.4 Machine Learning Tasks
 - 2.5 Linear Regression
 - 2.8.1 Gradient Descent
- 3 Neural Networks

1 Computational Mechanics Meets Artificial Intelligence & 2 Fundamental Concepts of Machine Learning

Leon Herrmann

Stefan Kollmannsberger

Chair of Data Engineering in Construction

Bauhaus-Universität Weimar

Paris, February 2025

*Deep Learning in Computational Mechanics – an introductory course,
Herrmann et al. 2025*



website



book

